

One role brief in. A full hiring package out.

A Claude plugin running eleven coordinated talent-acquisition agents. Tell it what you are hiring for. One research pass builds a per-role Context Pack from public data; eight experts write in parallel from the same pack; a deterministic renderer and a QA gate finish the job. Sourced, branded documents your team can use the same day.


ZERO SETUP COST Works with no connections. Public data plus a bundled baseline snapshot (BLS, Census, benchmark constants) is the default path. Research runs on a hard 12-call web budget.	EVERY NUMBER CITED Sourced (named public source + date) · Estimate (method shown) · Internal (your data). No receipt, no ship.
ASK SMALL OR LARGE "Write me a JD" runs the minimal chain. "Run talent one" runs all eight in parallel. Both pass the same gate.	UPDATES IN SECONDS Documents never pass through a model at render time. A comp change re-renders the kit in about one second , at zero AI cost.

THE EIGHT DOCUMENTS

01 · SENSEI Role Educational Brief Teaches a non-technical recruiter the domain: glossary, skills, career paths, say-this-not-that.	02 · JD-BOT Job Description Paste-ready posting with an annotation rail explaining why every line is written that way.	03 · ATLAS Talent Intelligence Map Supply, demand, geography, personas, comp. The research-heavy market read.	04 · HUNTER Sourcing Playbook Tiered boolean library, channel mix, timing windows, and honest pipeline math.
05 · SHAKESPEARE Outreach Campaign 8 touches over 21 days across email, LinkedIn, SMS, and phone. Zero recruiter cliches.	06 · INTERVIEW LAB Interview Plan Loop design, competency ownership grid, graded rubrics, and a debrief protocol.	07 · CALIBRATE HM Calibration Brief Market frame, comp anchor, question bank, and SLAs for the hiring-manager intake.	08 · RECRUITER SCREEN 30-Minute Screen Script A cold-runnable screen with timed runsheet and graded spot-check questions.

One research pass. Eight experts in parallel. One gate.

Scout researches the role once and writes a per-role **Context Pack**: one market read, one comp range, one persona set, with provenance on every number. Every artifact agent inherits the pack instead of re-researching, so every document agrees with every other. Agents emit content only; a deterministic renderer owns the HTML.



THE PACK COMPOUNDS	GAPS, NOT FILLER	CONNECTIONS OPTIONAL
The first run pays for the research. Same-role documents afterward take minutes; a comp update re-renders in about a second with zero AI spend.	A number with no receipt does not ship. Missing data is labeled as missing, never invented.	An ATS or your docs upgrade specific agents. Everything runs without them.

Install the file. Say “set up talent one.”

No files to edit, no configuration screens. A guided conversation captures your company once; every future run reuses it. Every step has a default and a skip. Fastest path: paste the bundled research prompt into any Claude chat and drop the JSON in — the research step then never runs.

1 Type your company name That is the whole job. Talent One researches what you do, size, locations, open roles, and recent news it can use as outreach proof. You correct anything wrong; only confirmed facts are stated as facts.	2 Pick what you hire for Roles found on your careers page arrive as one-tap options. This seeds defaults; it never limits what you can run later.	3 Constraints, in your words Work authorization, clearance, onsite, licensure. You write the exact sentence a candidate should see; it renders verbatim forever. Pick “none” and no document will ever hint at an eligibility gate.	4 Connect tools, or don’t ATS and document connections upgrade specific agents. Everything works without them.	5 Choose your look Light, dark, or dancheong heritage color, plus an optional brand color. You see a preview, not a hex code.
--	--	--	---	--

~30 min full 8-document kit, cold, including research and the QA gate (measured live)	minutes a single document on a role the pack already knows	~1 sec an update — comp change, new trigger — re-rendered at zero AI cost	100% of quantitative claims carry a provenance pill or they do not ship
---	--	---	---

THEN JUST ASK, IN PLAIN WORDS

- “write me a JD for a senior recruiter”
- “build a talent map for this role”
- “make an interview plan”
- “screen guide for the ops manager search”
- “run talent one for a Staff ML Engineer”
→ the full 8-document kit, researched, machine-verified, QA-gated, delivered as it completes

START HERE — RUNS IN CLAUDE CODE & COWORK ONLY

Install the `.plugin` file in **Claude Code** or **Cowork** (the Claude desktop app), then say “set up talent one.”

It will not run in the regular Claude chat window — chat can’t spawn the kit’s eleven sub-agents, so “set up talent one” does nothing there. Generated documents are yours to use, edit, and publish; the kit itself stays © Onepromptman and is not for resale. Documents are drafts for professional review — comp and eligibility language should be checked by whoever owns those decisions.

